

# index

October 23, 2020

```
[1]: from my_funct import confusion_matrix_info, load_clean_data, \
      ↪ create_plot_of_feature_importances
```

## 0.1 Data Cleaning

```
[2]: import pandas as pd
      import numpy as np
```

```
[3]: df = pd.read_csv('../data/churn_data.csv')
```

```
[4]: df.head()
```

```
[4]:  state  account length  area code phone number international plan \
0     KS           128    415     382-4657                no
1     OH           107    415     371-7191                no
2     NJ           137    415     358-1921                no
3     OH            84    408     375-9999                yes
4     OK            75    415     330-6626                yes

      voice mail plan  number vmail messages  total day minutes  total day calls \
0                yes                25          265.1          110
1                yes                26          161.6          123
2                no                 0          243.4          114
3                no                 0          299.4           71
4                no                 0          166.7          113

      total day charge  ...  total eve calls  total eve charge \
0          45.07  ...           99          16.78
1          27.47  ...          103          16.62
2          41.38  ...          110          10.30
3          50.90  ...           88           5.26
4          28.34  ...          122          12.61

      total night minutes  total night calls  total night charge \
0          244.7           91          11.01
1          254.4          103          11.45
2          162.6          104           7.32
3          196.9           89           8.86
```

4	186.9	121	8.41
---	-------	-----	------

	total intl minutes	total intl calls	total intl charge \
0	10.0	3	2.70
1	13.7	3	3.70
2	12.2	5	3.29
3	6.6	7	1.78
4	10.1	3	2.73

	customer service calls	churn
0	1	False
1	1	False
2	0	False
3	2	False
4	3	False

[5 rows x 21 columns]

[5]: *#we need to convert our yes/no and true/false into 0's and 1's*

```
churn_dict = {False: 0, True: 1}
yes_no_dict = {'no': 0, 'yes': 1}
df['churn'].replace(churn_dict, inplace=True)
df['international plan'].replace(yes_no_dict, inplace=True)
df['voice mail plan'].replace(yes_no_dict, inplace=True)
```

[6]: df.head()

	state	account length	area code	phone number	international plan \
0	KS	128	415	382-4657	0
1	OH	107	415	371-7191	0
2	NJ	137	415	358-1921	0
3	OH	84	408	375-9999	1
4	OK	75	415	330-6626	1

	voice mail plan	number vmail messages	total day minutes	total day calls \
0	1	25	265.1	110
1	1	26	161.6	123
2	0	0	243.4	114
3	0	0	299.4	71
4	0	0	166.7	113

	total day charge	...	total eve calls	total eve charge \
0	45.07	...	99	16.78
1	27.47	...	103	16.62
2	41.38	...	110	10.30
3	50.90	...	88	5.26

4	28.34	...	122	12.61
---	-------	-----	-----	-------

	total night minutes	total night calls	total night charge \
0	244.7	91	11.01
1	254.4	103	11.45
2	162.6	104	7.32
3	196.9	89	8.86
4	186.9	121	8.41

	total intl minutes	total intl calls	total intl charge \
0	10.0	3	2.70
1	13.7	3	3.70
2	12.2	5	3.29
3	6.6	7	1.78
4	10.1	3	2.73

	customer service calls	churn
0	1	0
1	1	0
2	0	0
3	2	0
4	3	0

[5 rows x 21 columns]

```
[7]: #pop off unneeded columns
```

```
df = df.drop(['phone number', 'area code'], axis=1)
df.head()
```

```
[7]: state account length international plan voice mail plan \
0 KS 128 0 1
1 OH 107 0 1
2 NJ 137 0 0
3 OH 84 1 0
4 OK 75 1 0
```

	number vmail messages	total day minutes	total day calls \
0	25	265.1	110
1	26	161.6	123
2	0	243.4	114
3	0	299.4	71
4	0	166.7	113

	total day charge	total eve minutes	total eve calls	total eve charge \
0	45.07	197.4	99	16.78

1	27.47	195.5	103	16.62
2	41.38	121.2	110	10.30
3	50.90	61.9	88	5.26
4	28.34	148.3	122	12.61

	total night minutes	total night calls	total night charge \
0	244.7	91	11.01
1	254.4	103	11.45
2	162.6	104	7.32
3	196.9	89	8.86
4	186.9	121	8.41

	total intl minutes	total intl calls	total intl charge \
0	10.0	3	2.70
1	13.7	3	3.70
2	12.2	5	3.29
3	6.6	7	1.78
4	10.1	3	2.73

	customer service calls	churn
0	1	0
1	1	0
2	0	0
3	2	0
4	3	0

```
[8]: #lets check our target variables value count for balance

df.churn.value_counts()

#looks like we have a large imbalance, this is something we can fix using SMOTE
```

```
[8]: 0    2850
     1     483
     Name: churn, dtype: int64
```

```
[9]: #let's now prepare our data for the train_test_split

X = df.drop('churn', axis=1)
y = df.churn
```

```
[10]: #we must import the proper packages to perform train_test_split

from sklearn.model_selection import train_test_split, cross_val_score

X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=2021,
↳test_size=0.20)
```

## 0.2 Making Pipelines

```
[11]: #let's create a pipeline to do all of our preprocessing for us

from imblearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import make_column_selector, make_column_transformer
from sklearn.model_selection import GridSearchCV
from imblearn.over_sampling import SMOTE

[12]: preprocessing = make_column_transformer((OneHotEncoder(),  
→make_column_selector(dtype_include=object)),  
→(StandardScaler(),  
→make_column_selector(dtype_include=np.number),  
→SMOTE()))

preprocessing

[12]: ColumnTransformer(transformers=[('onehotencoder', OneHotEncoder(),  
<sklearn.compose._column_transformer.make_column_selector object at  
0x00000243491FC9D0>),  
→('standardscaler', StandardScaler(),  
<sklearn.compose._column_transformer.make_column_selector object at  
0x00000243491FCB50>)])

[13]: #fit and transform our preprocessing pipeline to our training data

preprocessing.fit_transform(X_train)

[13]: <2666x68 sparse matrix of type '<class 'numpy.float64'>'
with 47988 stored elements in Compressed Sparse Row format>

[14]: #the next thing we'll do is make separate pipelines for each model we want to  
→test  
#each of these pipelines will contain our preprocessing pipeline

from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import ExtraTreesClassifier
from sklearn.neighbors import KNeighborsClassifier

dt_pipeline = make_pipeline(preprocessing,  
→DecisionTreeClassifier(random_state=2021))
rf_pipeline = make_pipeline(preprocessing,  
→RandomForestClassifier(random_state=2021))
```

```

lr_pipeline = make_pipeline(preprocessing,␣
    ↳ LogisticRegression(random_state=2021))
et_pipeline = make_pipeline(preprocessing,␣
    ↳ ExtraTreesClassifier(random_state=2021))
kn_pipeline = make_pipeline(preprocessing, KNeighborsClassifier())

```

The purpose for the creation of various pipelines is to find the best performing model using our training data that can then be used to perform reliably on unseen data. In this context, we can use our best model to predict churn patterns for SyriaTel.

### 0.3 Creating Our Param\_Grids

```

[15]: #we need to create different param_grids for each pipeline

dt_param_grid = {
    'decisiontreeclassifier__criterion': ['entropy', 'gini'],
    'decisiontreeclassifier__splitter': ['best', 'random'],
    'decisiontreeclassifier__max_depth': [2, 5, 10],
    'decisiontreeclassifier__max_features': ['auto', 'sqrt', 'log2'],
    'decisiontreeclassifier__class_weight': ['none', 'balanced']
}

rf_param_grid = {
    'randomforestclassifier__n_estimators': [100, 1000, 2000],
    'randomforestclassifier__max_depth': [2, 5, 10]
}

lr_param_grid = {
    'logisticregression__penalty': ['l1', 'l2', 'elasticnet', 'none'],
    'logisticregression__dual': [True, False],
    'logisticregression__solver': ['newton-cg', 'lbfgs', 'liblinear', 'sag',␣
    ↳ 'saga'],
    'logisticregression__multi_class': ['auto', 'ovr', 'multinomial'],
    'logisticregression__n_jobs': [10, 20, 30],
    'logisticregression__C': [0.01, 0.1, 0.5]
}

et_param_grid = {
    'extratreesclassifier__criterion': ['entropy', 'gini'],
    'extratreesclassifier__max_depth': [2, 5, 10],
    'extratreesclassifier__n_estimators': [100, 250, 500],
    'extratreesclassifier__max_features': ['auto', 'sqrt', 'log2'],
    'extratreesclassifier__class_weight': ['none', 'balanced']
}

```

```

}

kn_param_grid = {
    'kneighborsclassifier__n_neighbors': [5, 10, 20],
    'kneighborsclassifier__weights': ['uniform', 'distance'],
    'kneighborsclassifier__algorithm': ['auto', 'ball_tree', 'kd_tree',
    ↪ 'brute'],
    'kneighborsclassifier__leaf_size': [25, 50, 100],
    'kneighborsclassifier__p': [1, 2],
    'kneighborsclassifier__metric': ['minkowski', 'manhattan']
}

```

```

[16]: #we can now use each param_grid in our GridSearchCV alongside its coinciding
    ↪ model

```

## 0.4 DecisionTree

```

[17]: search_dt = GridSearchCV(dt_pipeline, dt_param_grid, n_jobs=-1)

search_dt.fit(X_train, y_train)

```

```

[17]: GridSearchCV(estimator=Pipeline(steps=[('columntransformer',
ColumnTransformer(transformers=[('onehotencoder',
OneHotEncoder(),
<sklearn.compose._column_transformer.make_column_selector object at
0x00000243491FC9D0>),
('standardscaler',
StandardScaler(),
<sklearn.compose._column_transformer.make_column_selector object at
0x00000243491FCB50>)])),
                    ('decisiontreeclassifier',
DecisionTreeClassifier(random_state=2021))]),
n_jobs=-1,
param_grid={'decisiontreeclassifier__class_weight': ['none',
'balanced'],
'decisiontreeclassifier__criterion': ['entropy',
'gini'],
'decisiontreeclassifier__max_depth': [2, 5, 10],
'decisiontreeclassifier__max_features': ['auto',
'sqrt',
'log2'],
'decisiontreeclassifier__splitter': ['best',
'random']})

```

```

[18]: #we can check out its best parameters

```

```
search_dt.best_params_
```

```
[18]: {'decisiontreeclassifier__class_weight': 'balanced',  
      'decisiontreeclassifier__criterion': 'entropy',  
      'decisiontreeclassifier__max_depth': 10,  
      'decisiontreeclassifier__max_features': 'auto',  
      'decisiontreeclassifier__splitter': 'best'}
```

```
[19]: #then assign that best model using best_estimator_ to a variable
```

```
best_dt_pipeline = search_dt.best_estimator_
```

```
[20]: #then check its f1 score using the training data
```

```
best_dt_cross_val = cross_val_score(best_dt_pipeline, X_train, y_train,  
    ↪scoring='f1')
```

## 0.5 RandomForest

```
[21]: search_rf = GridSearchCV(rf_pipeline, rf_param_grid, n_jobs=-1)
```

```
search_rf.fit(X_train, y_train)
```

```
[21]: GridSearchCV(estimator=Pipeline(steps=[('columntransformer',  
      ColumnTransformer(transformers=[('onehotencoder',  
      OneHotEncoder(),  
      <sklearn.compose._column_transformer.make_column_selector object at  
      0x00000243491FC9D0>),  
      ('standardscaler',  
      StandardScaler(),  
      <sklearn.compose._column_transformer.make_column_selector object at  
      0x00000243491FCB50>)])),  
      ('randomforestclassifier',  
      RandomForestClassifier(random_state=2021))]),  
      n_jobs=-1,  
      param_grid={'randomforestclassifier__max_depth': [2, 5, 10],  
                  'randomforestclassifier__n_estimators': [100, 1000,  
                                                             2000]})
```

```
[22]: #we can check out its best parameters
```

```
search_rf.best_params_
```

```
[22]: {'randomforestclassifier__max_depth': 10,  
      'randomforestclassifier__n_estimators': 2000}
```



```
[23]: #then assign that best model using best_estimator_ to a variable
```

```
best_rf_pipeline = search_rf.best_estimator_
```

```
[24]: #then check its f1 score using the training data
```

```
best_rf_cross_val = cross_val_score(best_rf_pipeline, X_train, y_train,
    ↪scoring='f1')
```

## 0.6 LogisticRegression

```
[25]: search_lr = GridSearchCV(lr_pipeline, lr_param_grid, n_jobs=-1)
```

```
search_lr.fit(X_train, y_train)
```

```
[25]: GridSearchCV(estimator=Pipeline(steps=[('columntransformer',
ColumnTransformer(transformers=[('onehotencoder',
OneHotEncoder(),
<sklearn.compose._column_transformer.make_column_selector object at
0x00000243491FC9D0>),
('standardscaler',
StandardScaler(),
<sklearn.compose._column_transformer.make_column_selector object at
0x00000243491FCB50>)])),
                    ('logisticregre...
LogisticRegression(random_state=2021))]),
        n_jobs=-1,
        param_grid={'logisticregression__C': [0.01, 0.1, 0.5],
                    'logisticregression__dual': [True, False],
                    'logisticregression__multi_class': ['auto', 'ovr',
                                                         'multinomial'],
                    'logisticregression__n_jobs': [10, 20, 30],
                    'logisticregression__penalty': ['l1', 'l2',
                                                    'elasticnet', 'none'],
                    'logisticregression__solver': ['newton-cg', 'lbfgs',
                                                    'liblinear', 'sag',
                                                    'saga']})
```

```
[26]: #we can check out its best parameters
```

```
search_lr.best_params_
```

```
[26]: {'logisticregression__C': 0.01,
      'logisticregression__dual': False,
      'logisticregression__multi_class': 'auto',
      'logisticregression__n_jobs': 10,
      'logisticregression__penalty': 'l2',
```

```
'logisticregression__solver': 'newton-cg'}
```

```
[27]: #then assign that best model using best_estimator_ to a variable
```

```
best_lr_pipeline = search_lr.best_estimator_
```

```
[28]: #then check its f1 score using the training data
```

```
best_lr_cross_val = cross_val_score(best_lr_pipeline, X_train, y_train,  
    ↪scoring='f1')
```

## 0.7 ExtraTrees

```
[29]: search_et = GridSearchCV(et_pipeline, et_param_grid, n_jobs=-1)
```

```
search_et.fit(X_train, y_train)
```

```
[29]: GridSearchCV(estimator=Pipeline(steps=[('columntransformer',  
ColumnTransformer(transformers=[('onehotencoder',  
OneHotEncoder(),  
<sklearn.compose._column_transformer.make_column_selector object at  
0x00000243491FC9D0>),  
('standardscaler',  
StandardScaler(),  
<sklearn.compose._column_transformer.make_column_selector object at  
0x00000243491FCB50>)])),  
                    ('extratreesclassifier',  
ExtraTreesClassifier(random_state=2021))]),  
                n_jobs=-1,  
                param_grid={'extratreesclassifier__class_weight': ['none',  
                                                                    'balanced'],  
                            'extratreesclassifier__criterion': ['entropy', 'gini'],  
                            'extratreesclassifier__max_depth': [2, 5, 10],  
                            'extratreesclassifier__max_features': ['auto', 'sqrt',  
                                                                    'log2'],  
                            'extratreesclassifier__n_estimators': [100, 250, 500]})
```

```
[30]: #we can check out its best parameters
```

```
search_et.best_params_
```

```
[30]: {'extratreesclassifier__class_weight': 'balanced',  
       'extratreesclassifier__criterion': 'gini',  
       'extratreesclassifier__max_depth': 10,  
       'extratreesclassifier__max_features': 'auto',  
       'extratreesclassifier__n_estimators': 500}
```

```
[31]: #then assign that best model using best_estimator_ to a variable
```

```
best_et_pipeline = search_et.best_estimator_
```

```
[32]: #then check its f1 score using the training data
```

```
best_et_cross_val = cross_val_score(best_et_pipeline, X_train, y_train,  
    ↪scoring='f1')
```

## 0.8 KNeighbors

```
[33]: search_kn = GridSearchCV(kn_pipeline, kn_param_grid, n_jobs=-1)
```

```
search_kn.fit(X_train, y_train)
```

```
[33]: GridSearchCV(estimator=Pipeline(steps=[('columntransformer',  
ColumnTransformer(transformers=[('onehotencoder',  
OneHotEncoder(),  
<sklearn.compose._column_transformer.make_column_selector object at  
0x00000243491FC9D0>),  
('standardscaler',  
StandardScaler(),  
<sklearn.compose._column_transformer.make_column_selector object at  
0x00000243491FCB50>)])),  
                    ('kneighborscla...  
                    KNeighborsClassifier())]),  
n_jobs=-1,  
param_grid={'kneighborsclassifier__algorithm': ['auto',  
                                                'ball_tree',  
                                                'kd_tree',  
                                                'brute'],  
            'kneighborsclassifier__leaf_size': [25, 50, 100],  
            'kneighborsclassifier__metric': ['minkowski',  
                                             'manhattan'],  
            'kneighborsclassifier__n_neighbors': [5, 10, 20],  
            'kneighborsclassifier__p': [1, 2],  
            'kneighborsclassifier__weights': ['uniform',  
                                             'distance']})
```

```
[34]: #we can check out its best parameters
```

```
search_kn.best_params_
```

```
[34]: {'kneighborsclassifier__algorithm': 'auto',  
      'kneighborsclassifier__leaf_size': 25,  
      'kneighborsclassifier__metric': 'minkowski',  
      'kneighborsclassifier__n_neighbors': 5,
```

```
'kneighborsclassifier__p': 1,
'kneighborsclassifier__weights': 'uniform'}
```

```
[35]: #then assign that best model using best_estimator_ to a variable
```

```
best_kn_pipeline = search_kn.best_estimator_
```

```
[36]: #then check its f1 score using the training data
```

```
best_kn_cross_val = cross_val_score(best_kn_pipeline, X_train, y_train,
↪scoring='f1')
```

## 0.9 Model F1 Score Means

```
[37]: print(f"RandomForest: {best_rf_cross_val.mean()}\n DecisionTree:
↪{best_dt_cross_val.mean()}\n KNeighbors: {best_kn_cross_val.mean()}\n
↪LogisticRegression: {best_lr_cross_val.mean()}\n ExtraTrees:
↪{best_et_cross_val.mean()}")
```

```
RandomForest: 0.6162381988178686
DecisionTree: 0.49161777125440886
KNeighbors: 0.48455377627943996
LogisticRegression: 0.2248114677359882
ExtraTrees: 0.6481831855744898
```

```
[38]: #based on this information, we chose to proceed with ExtraTrees
```

## 0.10 Final Model

```
[39]: #refit training data onto best model
```

```
best_et_pipeline.fit(X_train, y_train)
```

```
[39]: Pipeline(steps=[('columntransformer',
                        ColumnTransformer(transformers=[('onehotencoder',
                                                         OneHotEncoder(),
                                                         <sklearn.compose._column_transformer.make_column_selector object at
                                                         0x0000024343E2A160>),
                                                         ('standardscaler',
                                                         StandardScaler(),
                                                         <sklearn.compose._column_transformer.make_column_selector object at
                                                         0x000002434D65B790>)])),
                        ('extratreesclassifier',
                        ExtraTreesClassifier(class_weight='balanced', max_depth=10,
                                             n_estimators=500, random_state=2021))])
```

```
[40]: #checking the f1 score of best model using the training and testing data
#similar for both indicating that we have a reliable model
```

```

f1_score_train_mean = best_et_cross_val.mean()
f1_score_test_mean = cross_val_score(best_et_pipeline, X_test, y_test,
    ↳scoring='f1').mean()

print(f'Training F1 Score: {f1_score_train_mean}')
print(f'Testing F1 Score: {f1_score_test_mean}')

```

Training F1 Score: 0.6481831855744898

Testing F1 Score: 0.6179055216206919

```

[41]: from sklearn.metrics import accuracy_score

#checking the accuracy of best model using the training and testing data
#similar for both indicating that we have a reliable model

train_preds = best_et_pipeline.predict(X_train)
test_preds = best_et_pipeline.predict(X_test)

print(f'Training Accuracy: {accuracy_score(y_train, train_preds)}')
print(f'Testing Accuracy: {accuracy_score(y_test, test_preds)}')

```

Training Accuracy: 0.9253563390847712

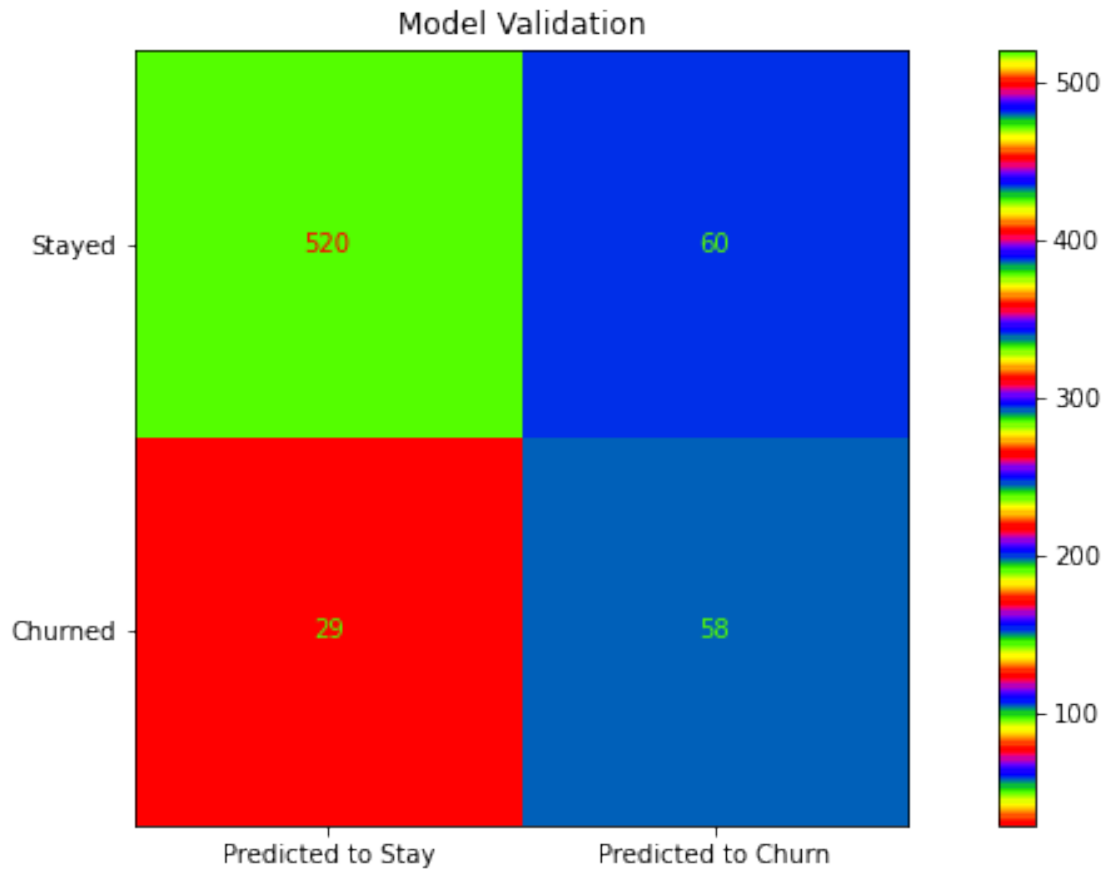
Testing Accuracy: 0.8665667166416792

## 0.11 Final Model Confusion Matrix

```

[42]: confusion_matrix_info(best_et_pipeline, X_test, y_test, save_path='images/
    ↳final_confusion_matrix.png')

```



[42]: (None, <Figure size 720x360 with 2 Axes>)

<Figure size 432x288 with 0 Axes>

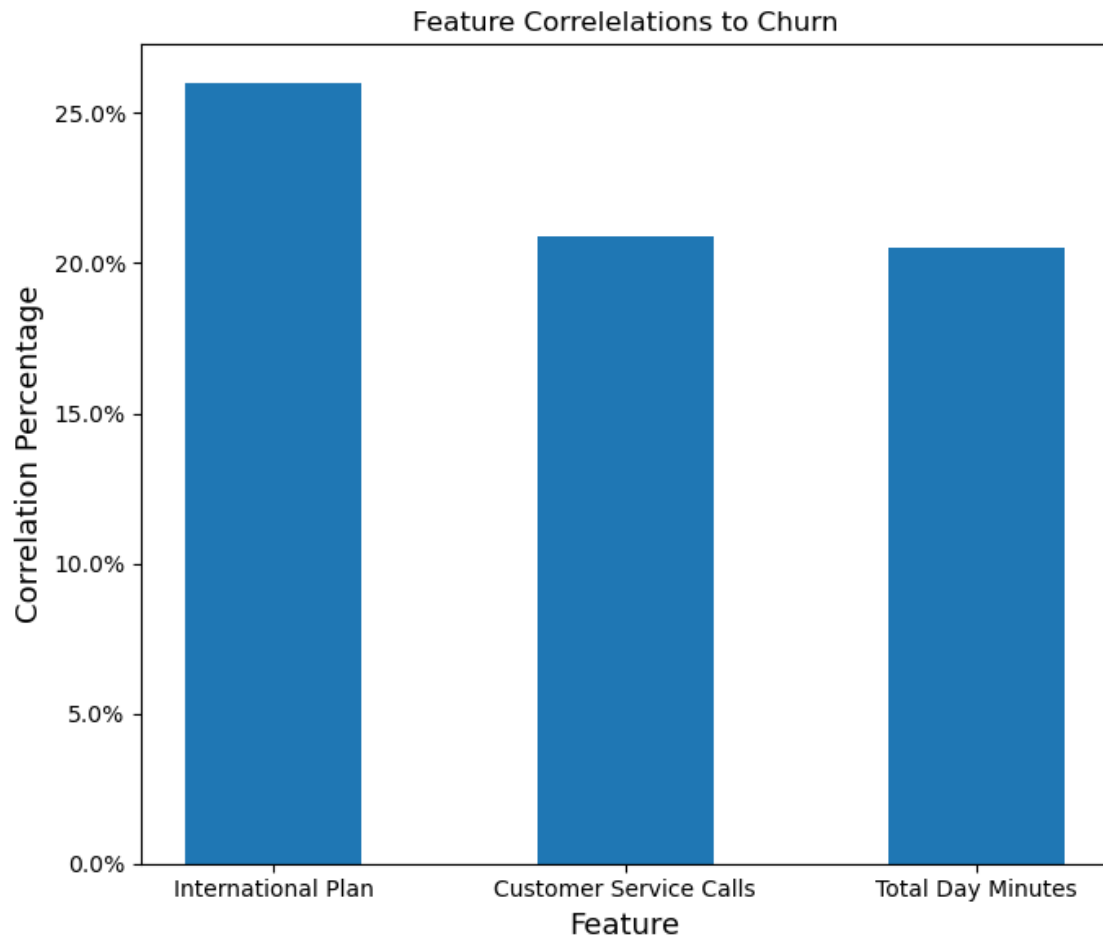
## 0.12 Other Visualizations

```
[43]: top_17 = [('total intl calls', 0.007662110949917248),
('total intl charge', 0.0033470091004430204),
('total day charge', 0.003259961859933837),
('total eve minutes', 0.003162403246382117),
('total night calls', 0.00277449674239552),
('account length', 0.002581535009061759),
('international plan', 0.002109226781666918),
('voice mail plan', 0.002103748904989105),
('total intl minutes', 0.0018711680184347753),
('total eve charge', 0.0017474977005077679),
('total day calls', 0.0017101284391014342),
('total night charge', 0.001615454532337387),
('number vmail messages', 0.0015464521017542058),
```

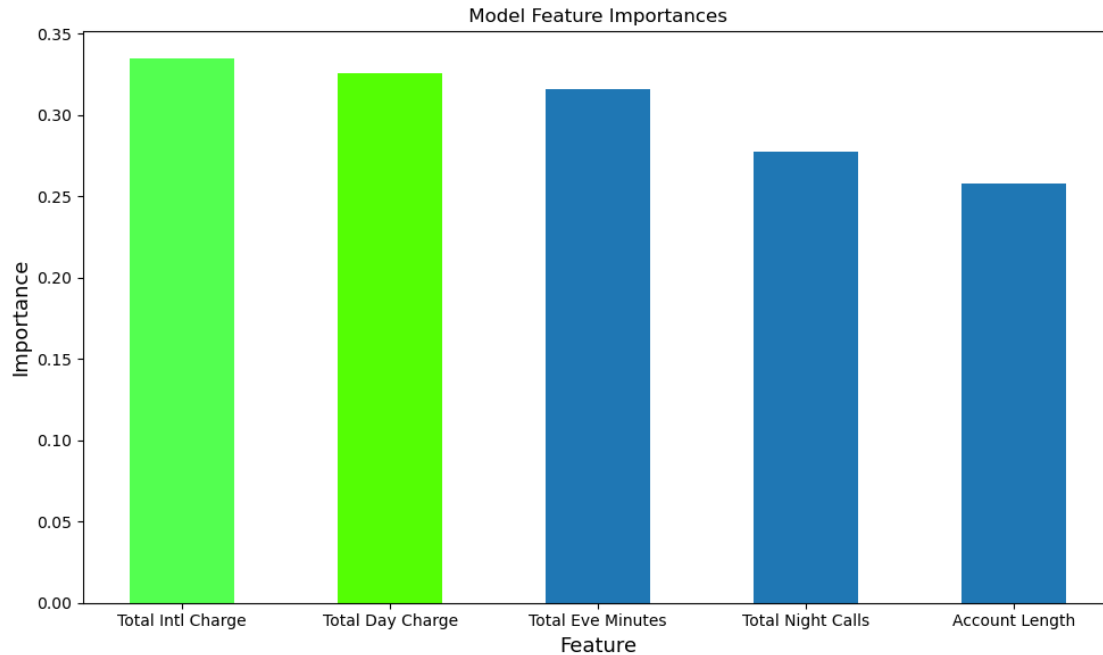
```
( 'customer service calls', 0.0012120735471571925),
( 'total eve calls', 0.001163035733859368),
( 'total night minutes', 0.000949022695754378),
( 'total day minutes', 0.0009380786304271744)]
```

```
[44]: min_percent = 0.12
corr = df.corr().churn
corr=dict(corr)
create_plot_of_feature_importances(corr, title='Feature Correlelations to
↳Churn', top_num=3, width=0.5, percent=True)
create_plot_of_feature_importances(dict(top_17), figsize=(10, 6),
↳bar_colors=['#53FF50', '#54FE04', '#1F77B4', '#1F77B4',
↳'#1F77B4', '#54FE00'], top_num=5, width=0.5, title='Model Feature
↳Importances', ylabel='Importance')

['International Plan', 'Customer Service Calls', 'Total Day Minutes']
[25.985184734548415, 20.874999878379207, 20.515082926138778]
['Total Intl Charge', 'Total Day Charge', 'Total Eve Minutes', 'Total Night
Calls', 'Account Length']
[0.334700910044302, 0.3259961859933837, 0.31624032463821167, 0.277449674239552,
0.2581535009061759]
```







```
[45]: train_f1_scores_dict = {
    "F1 Score": [
        cross_val_score(best_kn_pipeline, X_train, y_train, scoring='f1').
        ↪mean(),
        cross_val_score(best_rf_pipeline, X_train, y_train, scoring='f1').
        ↪mean(),
        cross_val_score(best_et_pipeline, X_train, y_train, scoring='f1').
        ↪mean(),
        cross_val_score(best_dt_pipeline, X_train, y_train, scoring='f1').mean()
    ]
}
train_f1_scores_df = pd.DataFrame(train_f1_scores_dict, index=["KNeighbors",
    ↪"RandomForest", "ExtraTrees", "DecisionTree"])
train_f1_scores_df
```

```
[45]:          F1 Score
KNeighbors    0.484554
RandomForest  0.616238
ExtraTrees    0.648183
DecisionTree  0.491618
```

```
[46]: test_f1_scores_dict = {
    "F1 Score": [
        cross_val_score(best_kn_pipeline, X_test, y_test, scoring='f1').mean(),
        cross_val_score(best_rf_pipeline, X_test, y_test, scoring='f1').mean(),
```

```

        cross_val_score(best_et_pipeline, X_test, y_test, scoring='f1').mean(),
        cross_val_score(best_dt_pipeline, X_test, y_test, scoring='f1').mean()
    ]
}
test_f1_scores_df = pd.DataFrame(test_f1_scores_dict, index=["KNeighbors",
↪ "RandomForest", "ExtraTrees", "DecisionTree"])
test_f1_scores_df

```

```

[46]:
      F1 Score
KNeighbors  0.355152
RandomForest 0.298128
ExtraTrees  0.617906
DecisionTree 0.414153

```